



**Trường Đại học Tôn Đức Thắng**

# Machine Learning Language Modeling

Giảng viên: Tiến sĩ Bùi Thanh Hùng  
Trưởng Lab Khoa học Phân tích dữ liệu và Trí tuệ nhân tạo  
Giám đốc chương trình Hệ thống thông tin  
Đại học Thủ Dầu Một  
Email: [tuhungphe@gmail.com](mailto:tuhungphe@gmail.com)  
Website: <https://sites.google.com/site/hungthanhbui1980/>



# Natural Language Processing

---

## Natural Language Processing

Cleanup, Tokenization

Stemming

Lemmatization

Part of Speech Tagging

Query Expansion

Parsing

Topic Segmentation and  
Recognition

Morphological Degmentation  
(Word/Sentences)

Information Retrieval and  
Extraction (IR)

Relationship Extraction

Named Entity Recognition  
(NER)

Sentiment Analysis/Sentence  
Boundary Disambiguation

Word sense and  
Disambiguation

Text Similarity

Coreference Resolution

Discourse Analysis

Machine Translation

Automatic Summarization/  
Paraphrasing

Natural Language Generation

Reasoning over  
Knowledge Based

Question Answering System

Dialog System

Image Captioning & other  
Multimodal Tasks

### Dependency Parsing

#### Raw sentence

He reckons the current account deficit will narrow to only 1.8 billion in September.



Part-of-speech tagging

#### POS-tagged sentence

He reckons the current account deficit will narrow to only 1.8 billion in September.


PRP VBZ DT JJ NN NN MD VB TO RB CD CD IN NNP .

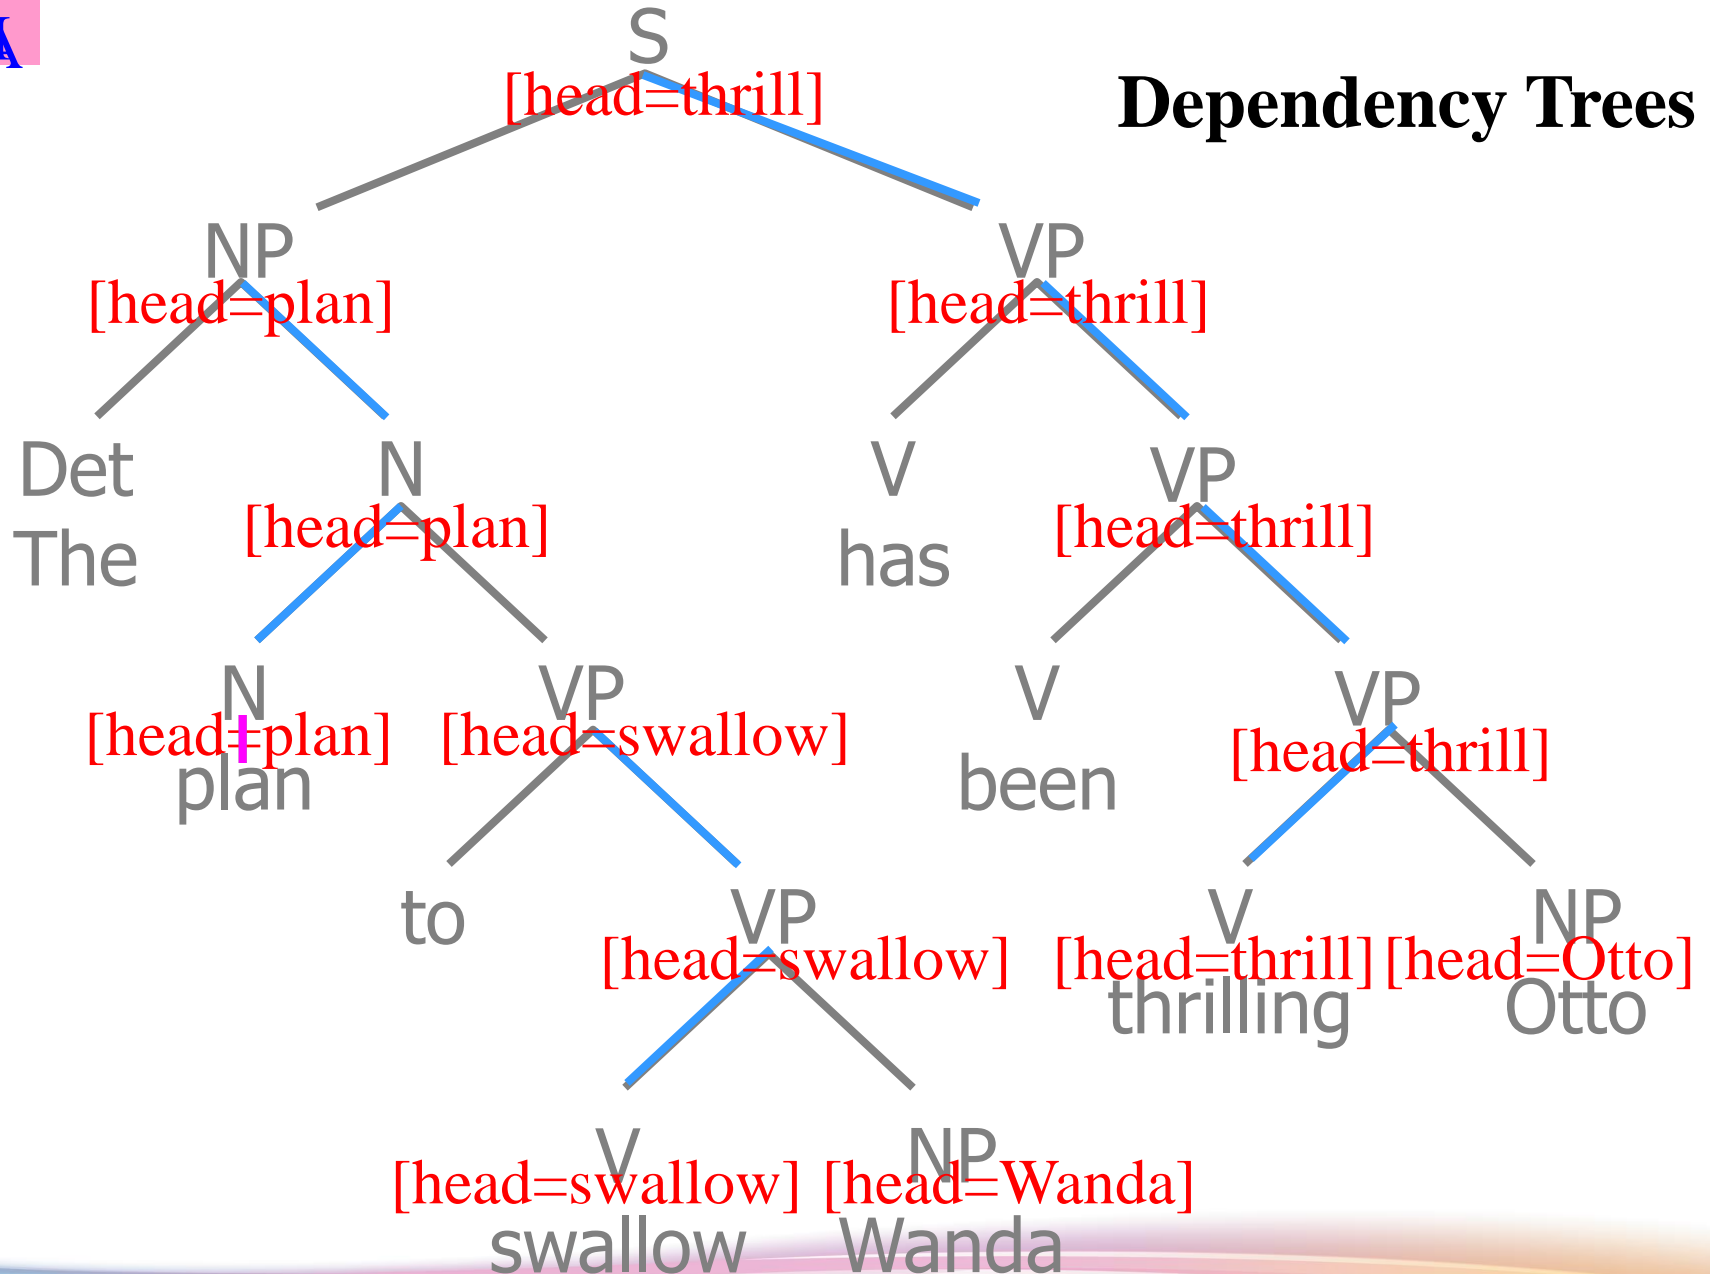


Word dependency parsing

#### Word dependency parsed sentence

He reckons the current account deficit will narrow to only 1.8 billion in September .





# Natural Language Processing

## Parsing (in Definite Clause Grammars)

s --> np, vp.

np --> det, noun.

np --> proper\_noun.

vp --> v, ng.

vp --> v.

det --> [a].      det --> [an].

det --> [the].

noun --> [apple].

noun --> [orange].


proper\_noun --> [john].

proper\_noun --> [mary].

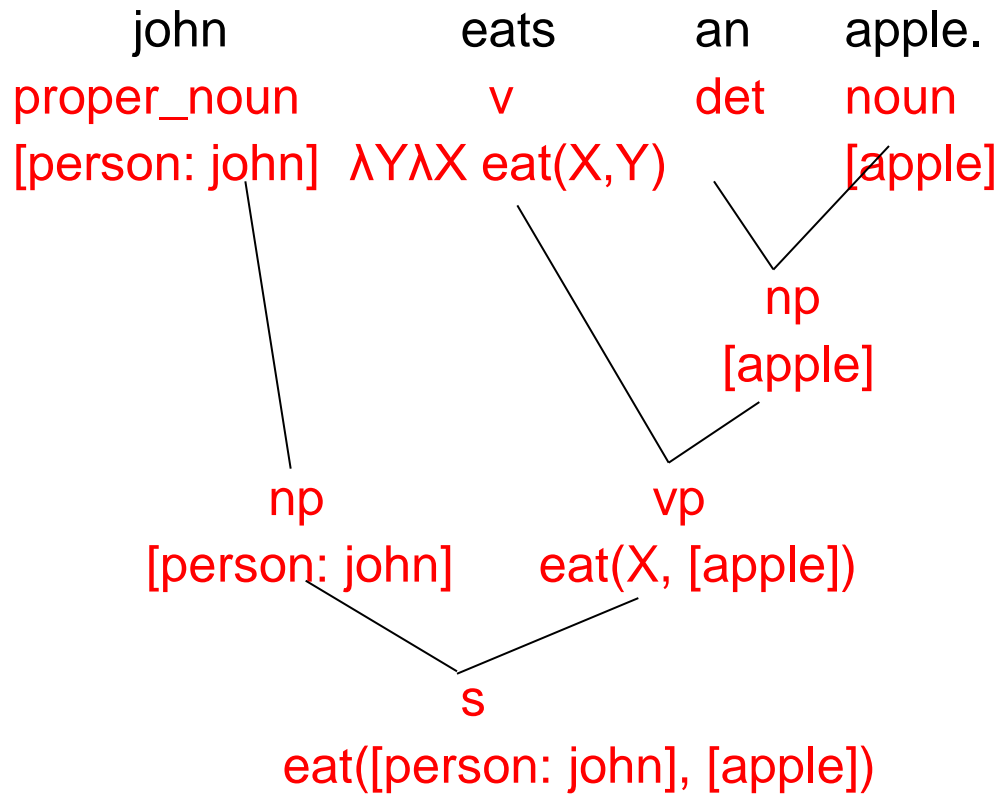
v --> [eats].

v --> [loves].


Eg.



## Semantic analysis



## Sem. Cat (Ontology)



## Bag of words - One hot vector

Nhà\_hàng này ngon, nhân\_viên nhiệt\_tình

→ [ 0, 0, 0, 0, 1, 1, 1, 1, 1, 0 ]

Đồ ăn ngon, nhưng mắc.

→ [ 1, 0, 1, 1, 0, 1, 0, 0, 0, 1 ]

Nhân\_viên chuyên\_nghiệp

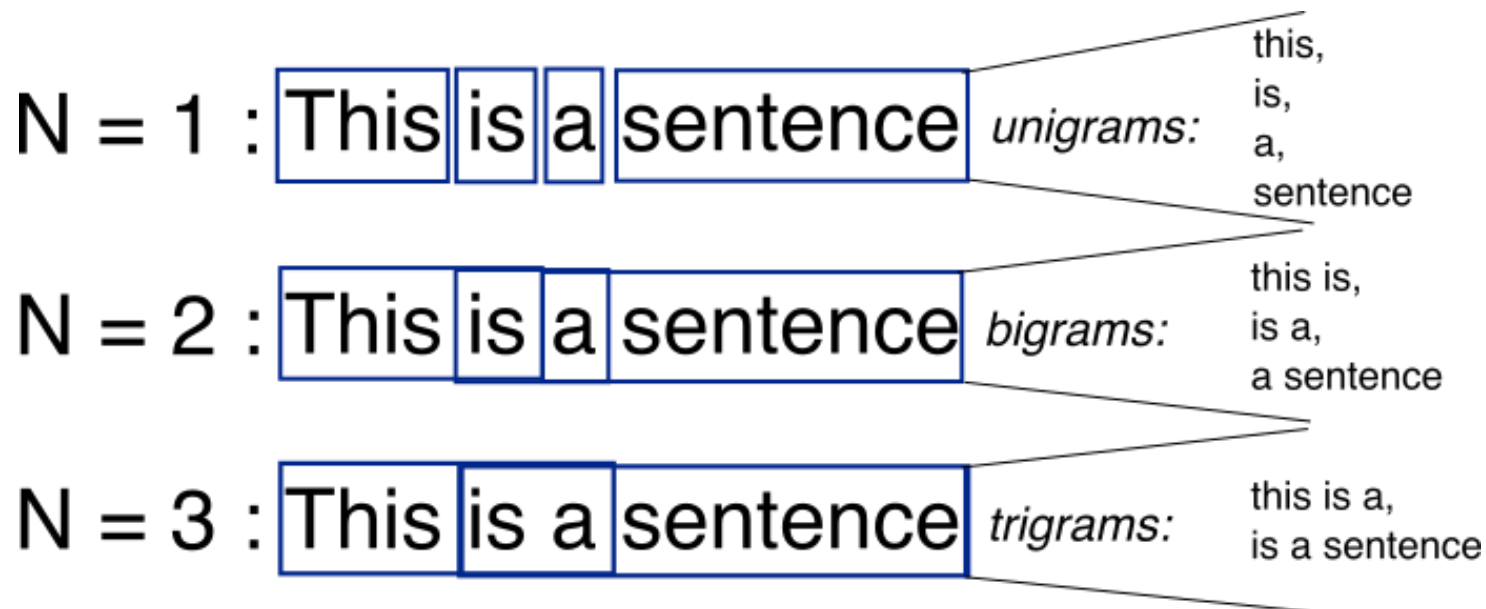
→ [ 0, 1, 0, 0, 0, 0, 0, 0, 1, 0 ]



0	ăn
1	chuyên_nghiệp
2	đồ
3	mắc
4	này
5	ngon
6	nhà_hàng
7	nhân_viên
8	nhiệt_tình
9	nhưng

## N-Gram

is a connected string of **N**







# Probabilistic Language Models

---

- Today's goal: assign a probability to a sentence
  - Machine Translation:
    - »  $P(\mathbf{high} \text{ winds tonite}) > P(\mathbf{large} \text{ winds tonite})$
  - Spell Correction
    - » The office is about fifteen **minuets** from my house
      - $P(\text{about fifteen } \mathbf{minutes} \text{ from}) > P(\text{about fifteen } \mathbf{minuets} \text{ from})$
  - Speech Recognition
    - »  $P(\text{I saw a van}) \gg P(\text{eyes awe of an})$
  - + Summarization, question-answering, etc., etc.!!

Why?

# Probabilistic Language Modeling

---

- Goal: compute the probability of a sentence or sequence of words:

$$P(W) = P(w_1, w_2, w_3, w_4, w_5 \dots w_n)$$

- Related task: probability of an upcoming word:

$$P(w_5 | w_1, w_2, w_3, w_4)$$

- A model that computes either of these:

$P(W)$  or  $P(w_n | w_1, w_2 \dots w_{n-1})$  is called a **language model**.

- Better: **the grammar** But **language model** or **LM** is standard

# How to compute $P(W)$

---

- How to compute this joint probability:
  - $P(\text{its, water, is, so, transparent, that})$
- Intuition: let's rely on the Chain Rule of Probability

# Reminder: The Chain Rule

---

- Recall the definition of conditional probabilities

$$p(B|A) = P(A,B)/P(A) \quad \text{Rewriting: } P(A,B) = P(A)P(B|A)$$

- More variables:

$$P(A,B,C,D) = P(A)P(B|A)P(C|A,B)P(D|A,B,C)$$

- The Chain Rule in General

$$P(x_1, x_2, x_3, \dots, x_n) = P(x_1)P(x_2|x_1)P(x_3|x_1, x_2) \dots P(x_n|x_1, \dots, x_{n-1})$$

# The Chain Rule applied to compute joint probability of words in sentence

---

$$P(w_1 w_2 \dots w_n) = \prod_i P(w_i \mid w_1 w_2 \dots w_{i-1})$$

$P(\text{"its water is so transparent"}) =$

$P(\text{its}) \times P(\text{water} \mid \text{its}) \times P(\text{is} \mid \text{its water})$

$\times P(\text{so} \mid \text{its water is}) \times P(\text{transparent} \mid \text{its water is so})$

# How to estimate these probabilities

---

- Could we just count and divide?

$$P(\text{the} \mid \text{its water is so transparent that}) = \frac{\text{Count}(\text{its water is so transparent that the})}{\text{Count}(\text{its water is so transparent that})}$$

- No! Too many possible sentences!
- We'll never see enough data for estimating these

# Markov Assumption

---



Andrei Markov

- Simplifying assumption:

$P(\text{the} \mid \text{its water is so transparent that}) \gg P(\text{the} \mid \text{that})$

- Or maybe

$P(\text{the} \mid \text{its water is so transparent that}) \gg P(\text{the} \mid \text{transparent that})$



# Markov Assumption

---

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i \mid w_{i-k} \dots w_{i-1})$$

- In other words, we approximate each component in the product

$$P(w_i \mid w_1 w_2 \dots w_{i-1}) \approx P(w_i \mid w_{i-k} \dots w_{i-1})$$

# Simplest case: Unigram model

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i)$$

Some automatically generated sentences from a unigram model

fifth, an, of, futures, the, an, incorporated, a, a, the,  
inflation, most, dollars, quarter, in, is, mass

thrift, did, eighty, said, hard, 'm, july, bullish

that, or, limited, the

# Bigram model

- Condition on the previous word:

$$P(w_i \mid w_1 w_2 \dots w_{i-1}) \approx P(w_i \mid w_{i-1})$$

texaco, rose, one, in, this, issue, is, pursuing, growth, in, a,  
boiler, house, said, mr., gurria, mexico, 's, motion, control,  
proposal, without, permission, from, five, hundred, fifty, five,  
yen

outside, new, car, parking, lot, of, the, agreement, reached

this, would, be, a, record, november

# N-gram models

---

- We can extend to trigrams, 4-grams, 5-grams
- In general this is an insufficient model of language
  - because language has **long-distance dependencies**:  
“The computer which I had just put into the machine room on the fifth floor crashed.”
- But we can often get away with N-gram models

# Estimating bigram probabilities

---

- The Maximum Likelihood Estimate

$$P(w_i \mid w_{i-1}) = \frac{\textit{count}(w_{i-1}, w_i)}{\textit{count}(w_{i-1})}$$

$$P(w_i \mid w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

## An example

---

$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

$$P(\text{I} | \text{<s>}) = \frac{2}{3} = .67$$

$$P(\text{Sam} | \text{<s>}) = \frac{1}{3} = .33$$

$$P(\text{am} | \text{I}) = \frac{2}{3} = .67$$

$$P(\text{</s>} | \text{Sam}) = \frac{1}{2} = 0.5$$

$$P(\text{Sam} | \text{am}) = \frac{1}{2} = .5$$

$$P(\text{do} | \text{I}) = \frac{1}{3} = .33$$

## More examples:

---

### Berkeley Restaurant Project sentences

- can you tell me about any good cantonese restaurants close by
- mid priced thai food is what i'm looking for
- tell me about chez panisse
- can you give me a listing of the kinds of food that are available
- i'm looking for a good place to eat breakfast
- when is caffe venezia open during the day

## Raw bigram counts

- Out of 9222 sentences

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0



# Raw bigram probabilities

- Normalize by unigrams:

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

- Result:

	i	want	to	eat	chinese	food	lunch	spend
i	0.002	0.33	0	0.0036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
to	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
eat	0	0	0.0027	0	0.021	0.0027	0.056	0
chinese	0.0063	0	0	0	0	0.52	0.0063	0
food	0.014	0	0.014	0	0.00092	0.0037	0	0
lunch	0.0059	0	0	0	0	0.0029	0	0
spend	0.0036	0	0.0036	0	0	0	0	0

# Bigram estimates of sentence probabilities

---

$$\begin{aligned} P(<s> \text{ I want english food } </s>) = \\ & P(\text{I} | <s>) \\ & \times P(\text{want} | \text{I}) \\ & \times P(\text{english} | \text{want}) \\ & \times P(\text{food} | \text{english}) \\ & \times P(</s> | \text{food}) \\ & = .000031 \end{aligned}$$

# What kinds of knowledge?

---

- $P(\text{english} | \text{want}) = .0011$
- $P(\text{chinese} | \text{want}) = .0065$
- $P(\text{to} | \text{want}) = .66$
- $P(\text{eat} | \text{to}) = .28$
- $P(\text{food} | \text{to}) = 0$
- $P(\text{want} | \text{spend}) = 0$
- $P(i | \langle s \rangle) = .25$

# Practical Issues

---

- We do everything in log space
  - Avoid underflow
  - (also adding is faster than multiplying)

$$\log(p_1 \cdot p_2 \cdot p_3 \cdot p_4) = \log p_1 + \log p_2 + \log p_3 + \log p_4$$

# Estimating N-gram Probabilities

---

- Estimating N-gram Probabilities
- Evaluation and Perplexity

# Evaluation: How good is our model?

---

- Does our language model prefer good sentences to bad ones?
  - Assign higher probability to “real” or “frequently observed” sentences
    - Than “ungrammatical” or “rarely observed” sentences?
- We train parameters of our model on a **training set**.
- We test the model’s performance on data we haven’t seen.
  - A **test set** is an unseen dataset that is different from our training set, totally unused.
  - An **evaluation metric** tells us how well our model does on the test set.

# Perplexity

The best language model is one that best predicts an unseen test set

- Gives the highest  $P(\text{sentence})$

Perplexity is the inverse probability of the test set, normalized by the number of words:

$$PP(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

Chain rule:

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$

For bigrams:

$$PP(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$

**Minimizing perplexity is the same as maximizing probability**

# Dict to Vec

---

```
measurements = [  
    {'city': 'Dubai', 'temperature': 33.},  
    {'city': 'London', 'temperature': 12.},  
    {'city': 'San Francisco', 'temperature': 18.},  
]  
  
from sklearn.feature_extraction import DictVectorizer  
vec = DictVectorizer()  
print(vec.fit_transform(measurements).toarray())
```



# Dict to Vec

```
measurements = [  
    {'city': 'Dubai', 'temperature': 33.},  
    {'city': 'London', 'temperature': 12.},  
    {'city': 'San Francisco', 'temperature': 18.},  
]  
  
from sklearn.feature_extraction import DictVectorizer  
vec = DictVectorizer()  
print(vec.fit_transform(measurements).toarray())
```

```
[[ 1.  0.  0. 33.]  
 [ 0.  1.  0. 12.]  
 [ 0.  0.  1. 18.]]
```

# Count Vectorize

---

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
```

# Count Vectorize

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
```

(0, 8)	1
(0, 3)	1
(0, 6)	1
(0, 2)	1
(0, 1)	1
(1, 8)	1
(1, 3)	1
(1, 6)	1
(1, 1)	1
(1, 5)	2
(2, 6)	1
(2, 0)	1
(2, 7)	1
(2, 4)	1
(3, 8)	1
(3, 3)	1
(3, 6)	1
(3, 2)	1
(3, 1)	1

# Count Vectorize

---

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
print(X)
print(vectorizer.get_feature_names())
```

# Count Vectorize

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
print(X)
print(vectorizer.get_feature_names())
```

(0, 8)	1
(0, 3)	1
(0, 6)	1
(0, 2)	1
(0, 1)	1
(1, 8)	1
(1, 3)	1
(1, 6)	1
(1, 1)	1
(1, 5)	2
(2, 6)	1
(2, 0)	1
(2, 7)	1
(2, 4)	1
(3, 8)	1
(3, 3)	1
(3, 6)	1
(3, 2)	1
(3, 1)	1

```
['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this']
```

# Count Vectorize

---

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
print(X)
print(vectorizer.get_feature_names())

print(vectorizer.vocabulary_.get('document'))
print(vectorizer.vocabulary_.get('and'))
```

# Count Vectorize

---

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
print(X)
print(vectorizer.get_feature_names())
```

```
vectorizer.transform(['Something completely new.']).toarray()
```

# Count Vectorize

```
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer()
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X = vectorizer.fit_transform(corpus)
print(X)
print(vectorizer.get_feature_names())
```

```
vectorizer.transform(['Something completely new.']).toarray()
```

```
array([[0, 0, 0, 0, 0, 0, 0, 0, 0]])
```



# Bigram

---

```
bigram_vectorizer = CountVectorizer(ngram_range=(1, 2),  
                                   token_pattern=r'\b\w+\b', min_df=1)  
analyze = bigram_vectorizer.build_analyzer()  
print(analyze('Bi-grams are cool!'))
```

# Bigram

---

```
bigram_vectorizer = CountVectorizer(ngram_range=(1, 2),  
                                   token_pattern=r'\b\w+\b', min_df=1)  
analyze = bigram_vectorizer.build_analyzer()  
print(analyze('Bi-grams are cool!'))
```

```
['bi', 'grams', 'are', 'cool', 'bi grams', 'grams are', 'are cool']
```

# Bigram

---

```
from sklearn.feature_extraction.text import CountVectorizer
bigram_vectorizer = CountVectorizer(ngram_range=(1, 2),
                                   token_pattern=r'\b\w+\b', min_df=1)
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X_2 = bigram_vectorizer.fit_transform(corpus).toarray()
print(X_2)
```

# Bigram

---

```
from sklearn.feature_extraction.text import CountVectorizer
bigram_vectorizer = CountVectorizer(ngram_range=(1, 2),
                                   token_pattern=r'\b\w+\b', min_df=1)
corpus = [
    'This is the first document.',
    'This is the second second document.',
    'And the third one.',
    'Is this the first document?',
]
X_2 = bigram_vectorizer.fit_transform(corpus).toarray()
print(X_2)
```

```
[[0 0 1 1 1 1 1 0 0 0 0 0 1 1 0 0 0 0 1 1 0]
 [0 0 1 0 0 1 1 0 0 2 1 1 1 0 1 0 0 0 1 1 0]
 [1 1 0 0 0 0 0 0 1 0 0 0 1 0 0 1 1 1 0 0 0]
 [0 0 1 1 1 1 0 1 0 0 0 0 1 1 0 0 0 0 1 0 1]]
```

# TF-IDF

---

- In a large text corpus, some words will be very present (e.g. “the”, “a”, “is” in English) hence carrying very little meaningful information about the actual contents of the document. If we were to feed the direct count data directly to a classifier those very frequent terms would shadow the frequencies of rarer yet more interesting terms.
- In order to re-weight the count features into floating point values suitable for usage by a classifier it is very common to use the tf-idf transform.

# TF-IDF

---

- $t$  — term (word)
- $d$  — document (set of words)
- $N$  — count of corpus
- corpus — the total document set

# TF-IDF

---

- Tf means **term-frequency** while tf-idf means term-frequency times **inverse document-frequency**:

$$\text{tf-idf}(t, d) = \text{tf}(t, d) * \log(N/(df + 1))$$

- $tf(t, d) = \text{count of } t \text{ in } d / \text{number of words in } d$
- $df(t) = \text{occurrence of } t \text{ in documents}$
- $idf(t) = \log(N/(df + 1))$

# TF-IDF

---

- Tf means **term-frequency** while tf-idf means term-frequency times **inverse document-frequency**:

$$\text{tf-idf}(t, d) = \text{tf}(t, d) * \log(N/(df + 1))$$

- The resulting tf-idf vectors are then normalized by the Euclidean norm

$$v_{norm} = \frac{v}{||v||_2} = \frac{v}{\sqrt{v_1^2 + v_2^2 + \dots + v_n^2}}$$



# TF-IDF

---

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

# TF-IDF

---

```
from sklearn.feature_extraction.text import TfidfTransformer
transformer = TfidfTransformer(smooth_idf=False)
counts = [[3, 0, 1],
          [2, 0, 0],
          [3, 0, 0],
          [4, 0, 0],
          [3, 2, 0],
          [3, 0, 2]]

tfidf = transformer.fit_transform(counts)
print(tfidf.toarray())
```

# TF-IDF

---

```
from sklearn.feature_extraction.text import TfidfTransformer
transformer = TfidfTransformer(smooth_idf=False)
counts = [[3, 0, 1],
          [2, 0, 0],
          [3, 0, 0],
          [4, 0, 0],
          [3, 2, 0],
          [3, 0, 2]]

tfidf = transformer.fit_transform(counts)
print(tfidf.toarray())
```

```
[[0.81940995 0.          0.57320793]
 [1.          0.          0.          ]
 [1.          0.          0.          ]
 [1.          0.          0.          ]
 [0.47330339 0.88089948 0.          ]
 [0.58149261 0.          0.81355169]]
```

# TF-IDF

---

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

Term:

Document:

Df:

Tf:

Tf-idf:

# TF-IDF

---

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

Term: 3

Document: 6

$tf(t,d) = \text{count of } t \text{ in } d / \text{number of words in } d$

$df(t) = \text{occurrence of } t \text{ in documents}$

$idf(t) = \log(N / (df + 1))$

$tf-idf(t, d) = tf(t, d) * \log(N / (df + 1))$

# TF-IDF

---

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

$$n = 6$$

$$\text{df}(t)_{\text{term1}} = 6$$

$$\text{idf}(t)_{\text{term1}} = \log \frac{n}{\text{df}(t)} + 1 = \log(1) + 1 = 1$$

$$\text{tf-idf}_{\text{term1}} = \text{tf} \times \text{idf} = 3 \times 1 = 3$$

# TF-IDF

---

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

Term:

Document:

Df(term 1):

Tf(term 1):

Tf-idf(term 1):

Term:

Document:

Df(term 2):

Tf(term 2):

Tf-idf(term 2):

Term:

Document:

Df(term 3):

Tf(term 3):

Tf-idf(term 3):

# TF-IDF

```
counts = [[3, 0, 1],  
          [2, 0, 0],  
          [3, 0, 0],  
          [4, 0, 0],  
          [3, 2, 0],  
          [3, 0, 2]]
```

$$n = 6$$

$$\text{df}(t)_{\text{term1}} = 6$$

$$\text{idf}(t)_{\text{term1}} = \log \frac{n}{\text{df}(t)} + 1 = \log(1) + 1 = 1$$

$$\text{tf-idf}_{\text{term1}} = \text{tf} \times \text{idf} = 3 \times 1 = 3$$

$$\text{tf-idf}_{\text{term2}} = 0 \times (\log(6/1) + 1) = 0$$

$$\text{tf-idf}_{\text{term3}} = 1 \times (\log(6/2) + 1) \approx 2.0986$$

$$\text{tf-idf}_{\text{raw}} = [3, 0, 2.0986]$$

$$\frac{[3, 0, 2.0986]}{\sqrt{(3^2 + 0^2 + 2.0986^2)}} = [0.819, 0, 0.573]$$



# TF-IDF

---

Furthermore, the default parameter `smooth_idf=True` adds “1” to the numerator and denominator as if an extra document was seen containing every term in the collection exactly once, which prevents zero divisions:

# TF-IDF

---

## TF-IDF

- TF-IDF stands for term frequency-inverse document frequency.
- TF-IDF is used to evaluate how important a word is to a document in a collection or corpus.


**TF: Term Frequency**, the number of times a term occurs in a document.

**IDF: Inverse Document Frequency**, which measures the important of a term.


$$\text{IDF} = \log \frac{1+n}{1+DF} + 1$$

$$\text{TF-IDF} = \text{TF} \times \text{IDF}$$

## Word Embeddings



Male-Female



Verb tense

## Word Embeddings


### CBOW, Skip-gram

#### - Continuous Bag of Words (CBOW) (Tomas Mikolov et al., 2013):

This model predicts from the present context-based in large corpus volumes with a sliding window

Ex: “Con mèo ngồi trên sàn”

Sliding window = 2



## Word Embeddings

**CBOW** (Tomas Mikolov et al., 2013, 16840 cited – 15/07/2020)

**Skip-gram** (Tomas Mikolov et al., 2013, 16840 cited – 15/07/2020)

**GloVe** (Stanford NLP Research, 2014, 15159 cited – 15/07/2020)

**FastText** (Facebook AI Research 2016, 5773 cited – 15/07/2020)

**BERT** (Google AI, 2018: 7529 cited – 15/07/2020)

(Pre-training of Deep Bidirectional Transformers for Language Understanding)

Decorative wavy lines in shades of blue, purple, and orange at the bottom of the slide.

# Mid Term

---

- Crawl data from vnexpress (at least 5 topic, 40 documents/1 topic)
- Output:
  - Pre processing (stopword, html,tokenize...)
  - Extract TF-IDF feature
  - Using SVM for training of document classification
  - Evaluate by Accuracy